

CFD Lab – Lecture 3

Message Passing Interface - MPI

Introduction

Message Passing Interface (MPI) is a library that allows communication between different processes

- Enables explicit passing of the data between processes.
- From several cores on a single CPU till thousands of cores on SuperMUC

Hallo World in MPI

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char* argv[])
{
    int my_id;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_id);

    printf("Hello from rank %d.\n", my_id);
    MPI_Finalize();
}
```

1. Include `<mpi.h>` header file
2. Initialize MPI
3. Assign **IDs** to each process using `MPI_Comm_rank`
4. Each process prints its **ID**
5. Finalize MPI

Output of the program

```
C:\Users\stefa\source\repos\MPI_Test\x64\Release  
λ mpicc main.c -o MPI_Test
```

```
C:\Users\stefa\source\repos\MPI_Test\x64\Release  
λ mpiexec -n 4 MPI_Test.exe  
Hello from rank 2  
Hello from rank 3  
Hello from rank 1  
Hello from rank 0
```

- Code is compiled using **mpicc** flag
- In order to run our program we need to use **mpiexec** command
- Order of execution is not ensured
- As soon as process get assigned hardware resources it will perform its task


```
#include <stdio.h>
#include <mpi.h>
main(int argc, char* argv[])
{
    int my_id;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD,
    &my_id);

    printf("Hello from rank %d.\n", my_id);
    MPI_Finalize();
}
```

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char* argv[])
{
    int my_id;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD,
    &my_id);

    printf("Hello from rank %d.\n", my_id);
    MPI_Finalize();
}
```

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char* argv[])
{
    int my_id;

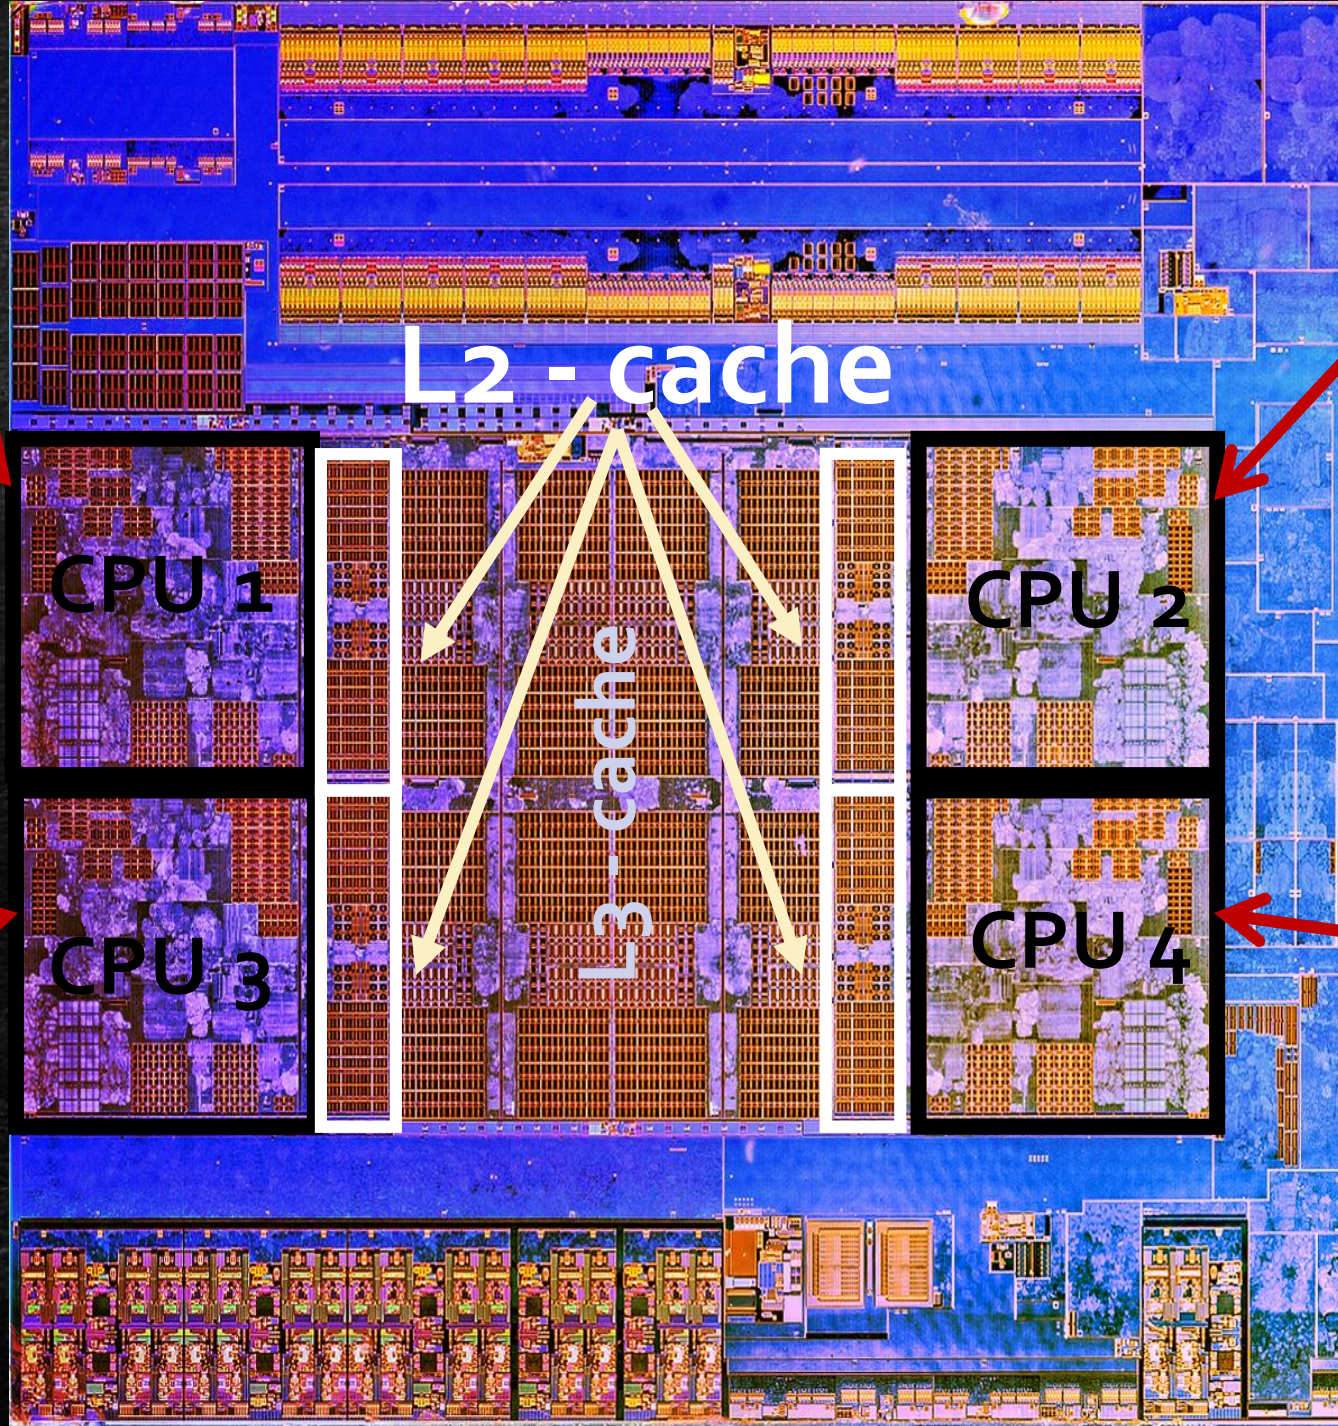
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD,
    &my_id);

    printf("Hello from rank %d.\n", my_id);
    MPI_Finalize();
}
```

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char* argv[])
{
    int my_id;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD,
    &my_id);

    printf("Hello from rank %d.\n", my_id);
    MPI_Finalize();
}
```



What to do when code should differ in each process?

- Each process has its own unique ID which help to differentiate between them
- For process dependent tasks we can use branching

```
if ( my_id = 0 )  
    // Routine_Calculate A  
else if ( my_id = 1 )  
    // Routine_Calculate B  
else if ( my_id = 2 )  
    // Routine_Calculate C
```


Sending and Receiving Messages

- Hello World example is easy, but there is no message exchange yet.
- Simplest Message Passing program can be implemented for only two processes (P_0 and P_1)
- Goal is to send number 23 from P_1 to P_0

MPI_Send(...)

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype,  
             int dest, int tag, MPI_Comm comm)
```

Input Parameters

buf - initial address of send buffer (choice)

count - number of elements in send buffer (nonnegative integer)

datatype - datatype of each send buffer element (handle)

dest - id of destination process (integer)

tag - message tag (integer)

comm - communicator (handle)

MPI_Recv(...)

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,  
             int source, int tag, MPI_Comm comm, MPI_Status *status)
```

Input Parameters

count - maximum number of elements in receive buffer (integer)

datatype - datatype of each receive buffer element (handle)

source - rank of source (integer)

tag - message tag (integer)

comm - communicator (handle)

Output Parameters

buf - initial address of receive buffer (choice)

status - status object (Status)

Send-Receive Example Code

```
#include <stdio.h>
#include <mpi.h>
void main(int argc, char* argv[])
{
    int my_id, number_to_receive, number_to_send = 23;
    MPI_Status status;
    MPI_Init (&argc, &argv);
    MPI_Comm_rank( MPI_COMM_WORLD, &my_id );

    if ( my_id == 0 )
    {
        MPI_Recv( &number_to_receive, 1, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
        printf( "Number received is: %d\n", number_to_receive );
    }
    else
    MPI_Send( &number_to_send, 1, MPI_INT, 0, 10, MPI_COMM_WORLD );

    MPI_Finalize();
}
```


Pitfalls

- Previously used Receive function call is so called blocking call
 - Receive will wait until message matching its requirements has been detected
- Send functions will try not to block, but this behavior is not guaranteed
 - With the growing size of data that is sent, this blocking behavior emerges
- For every `MPI_Send()` there must be pairing `MPI_Recv()`
- Otherwise Deadlock can occur

Deadlock Example

- Process with `id = 0` will wait to match receive from process `id = 1`
- Process with `id = 1` will wait to match receive from process `id = 0`
- Both processes blocked
- By switching statements for one of the processes problem is resolved

```
MPI_Comm_rank(comm, &my_id);

if (my_id == 0) {
    MPI_Send(sendbuf, count, MPI_INT, 1, tag, comm);
    MPI_Recv(recvbuf, count, MPI_INT, 1, tag, comm, &status);
} else if (my_id == 1) {
    MPI_Send(sendbuf, count, MPI_INT, 0, tag, comm);
    MPI_Recv(recvbuf, count, MPI_INT, 0, tag, comm, &status);
}
```


Useful Functions

```
int MPI_Allreduce(const void *sendbuf, void *recvbuf, int count,  
                  MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

```
int MPI_Reduce(const void *sendbuf, void *recvbuf, int count,  
               MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)
```

```
int MPI_Bcast( void *buffer, int count, MPI_Datatype datatype, int root,  
               MPI_Comm comm )
```


Synchronization helps to align all the processes

- In order to align all the processes at a certain point MPI_Barrier() can be used
- Hallo World with ordered output

```
C:\Users\stefa\source\repos\MPI_Tes
λ mpiexec.exe -n 4 MPI_Test.exe
Hello from rank 0
Hello from rank 1
Hello from rank 2
Hello from rank 3
```

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char* argv[])
{
    int my_id;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_id);

    for (int index = 0; index < 4; index++)
    {
        MPI_Barrier(MPI_COMM_WORLD);
        if (index == rank)
            printf("Hello from rank %d\n", my_id);
    }
    MPI_Finalize();
}
```